Pypomp: Inference for partially observed Markov process models in Python with JAX [DRAFT]

Aaron J. Abkemeier[1] AARONABK@UMICH.EDU
Jun Chen[2] CHENJUNC@UMICH.EDU
Kevin Tan[3] KEVTAN@WHARTON.UPENN.EDU
Jesse Wheeler[4] JESSE.WHEELER@USU.EDU
Bo Yang[5] YBB@UMICH.EDU
Kunyang He[1] KUNYANGHE@UMICH.EDU
Jonathan Terhorst[1] JONTH@UMICH.EDU
Aaron A. King[6] KINGAA@UMICH.EDU
Giles Hooker[3] GHOOKER@WHARTON.UPENN.EDU
Edward L. Ionides[1] IONIDES@UMICH.EDU

[1] *University of Michigan, Department of Statistics;* [2] *Duke University, Department of Statistics;* [3] *University of Pennsylvania, Department of Statistics;* [4] *Utah State University, Department of Mathematics & Statistics;* [5] *Columbia University, Department of Marketing;* [6] *University of Michigan, Department of Ecology & Evolutionary Biology*

Editor:

Abstract

Nonlinear, non-Gaussian partially observed Markov process (POMP) models arise in various areas including epidemiology, ecology, and robotics. Model-based inferences typically involve computationally intensive Monte Carlo methods. We introduce Pypomp, a Python library for data analysis via POMP models that uses JAX to take advantage of GPU hardware, automatic differentiation, and just-in-time compilation. By design, Pypomp supports model development and data analysis as well as method development, following the approach of the pomp R package. However, Pypomp obtains considerably higher performance: Pypomp running on one GPU can obtain the throughput of 8,000 CPU cores running pomp. A central algorithm for Pypomp is the particle filter, and Pypomp is faster on our benchmark tests than other GPU-compatible particle filter libraries. The particle filter is differentiable, with a tunable bias-variance tradeoff.

Keywords: pomp, JAX, GPU acceleration, likelihood-based inference

1 Introduction

Partially observed Markov process (POMP) models provide a flexible framework for modeling dynamic systems whose underlying state is not directly observed. POMP models are also known as state-space models or hidden Markov models, though these alternative names are sometimes restricted to continuous and discrete states respectively. Pypomp is a Python library combining advances in hardware, software, and algorithms to provide novel capabilities for POMP models, with an emphasis on inference via particle filters. The particle filter (PF), also known as sequential Monte Carlo, provides an estimate of the likelihood of

a POMP model (Arulampalam et al., 2002). The basic PF algorithm only needs to simulate from the process model, not evaluate its transition probabilities, a property known as plug-and-play (P&P) (Bretó et al., 2009). Pypomp focuses on P&P methods because they facilitate scientific modeling, enabling applications in epidemiology (Mietchen et al., 2024; Fox et al., 2022; Wen et al., 2024), ecology (Auger-Méthé et al., 2021; Marino et al., 2019; Blackwood et al., 2013), finance (Bretó, 2014), robotics (Jonschkowski et al., 2018) and other domains where nonlinear mechanistic stochastic dynamic models are pertinent. Pypomp provides P&P inference methods extending PF: maximum likelihood parameter estimation by iterated filtering (IF2) (Ionides et al., 2015), Bayesian inference by particle Markov chain Monte Carlo (Andrieu et al., 2010), and stochastic gradient descent via a differentiable particle filter (DPF) algorithm called a measurement-off-parameter (MOP) particle filter (Tan et al., 2026).

Pypomp follows the design principles of the R package pomp (King et al., 2016), providing a general framework for specifying models and inference methods, a range of pre-built models and methods, and various graphical and diagnostic tools to assist practical data analysis. Thus, Pypomp supports complete data analysis workflows, facilitating the evaluation of a range of models and statistical methods for the data under investigation. Models and methods have an interface designed so that new models can readily be tested on a range of pre-existing methods, and new methods can be tested on a range of pre-existing models. Although the method interface of Pypomp (and pomp) is not limited to P&P algorithms, prior experience with pomp has shown that P&P approaches are convenient for scientific data analysis. The computability of transition densities might in principle lead to non-P&P methods that are superior to P&P alternatives, though theoretical obstacles arise that have limited the practical scope of such algorithms (Roberts and Stramer, 2001; Karppinen et al., 2024).

Pypomp gains computationally over pomp by leveraging JAX (Bradbury et al., 2018) which provides just-in-time (JIT) compilation, automatic differentiation (AD), and capability to run on massively parallel hardware such as GPUs. When the state space is small and discrete, or the model is linear Gaussian or close enough for effective approximation using Gaussian-inspired filters such as the extended Kalman filter, the particle filter is unnecessary, and the Dynamax library (Linderman et al., 2025) provides JAX-based solutions. Other JAX-based libraries providing PF and DPF methods for POMP models include PF-JAX (Lysy et al., 2022) and cuthbert (Duffield, 2026). PyDPF (Brady et al., 2026) provides a DPF using Pytorch and focused on POMP models parameterized using neural networks. The Python library particles (Chopin and Papaspiliopoulos, 2020) provides a general framework for sequential Monte Carlo computation without AD. Pypomp distinguishes itself from these other packages by (ii) following the design decisions of pomp that support a full data analysis workflow; (iii) superior performance on our benchmark tests in Table 1. By design, Pypomp is not limited to PF methods, for example, the approximate Bayesian computing (ABC) P&P algorithm of Marjoram et al. (2003) has been implemented. However, PF has proved an effective tool for the development and analysis of POMP models in scientific contexts, and hence it is an important algorithm for Pypomp.

With enough particles and method replications to saturate it, one GPU runs iterated filtering as fast as 3,239 to 8,659 CPU cores running pomp (Appendix B). Such replications are routine in practice: likelihood maximization involves searching from many initial

parameter vectors, and Monte Carlo adjusted profile (MCAP) confidence intervals (Ionides et al., 2017), implemented in Pypomp as the `mcap` method, require many iterated filtering searches at each point of a likelihood profile. Profile confidence intervals can be superior to the widely used observed Fisher information construction, even in the case of linear Gaussian POMP models for which the likelihood and its gradients can be computed without Monte Carlo error (Wheeler and Ionides, 2025). MCAP algorithms can lead to efficient inference even in a limit where the Monte Carlo error grows as the amount of data increases (Ning et al., 2021). GPU hardware makes P&P likelihood-based inference feasible at previously intractable scales: one GPU fits a panel POMP model to the England and Wales measles data of Korevaar et al. (2020), over one million observations from 1,422 towns and cities, in 38 hours, compared with an estimated 4.8 months for `panelPomp` on a 36-core node (Section 3).

Algorithms to take advantage of JIT, AD, and GPU are an area of active development. Pypomp implements the MOP P&P DPF algorithm of Tan et al. (2026), with the resulting derivative being called DMOP and implemented as the `dmop` method. DMOP was demonstrated to be competitive on benchmark problems when combined with Adam stochastic gradient descent. A warm start using IF2 was found to be beneficial. Numerical random variable generation is a core operation for P&P methodology, and the development of sampling methods with favorable performance on GPU architectures is another current research topic. Pypomp includes optimized inverse transform sampling methods as alternatives to some `jax.random` samplers, as the rejection sampling methods used by `jax.random` are inefficient on GPUs (Giles, 2016; Giles and Beentjes, 2024; Temme, 1992). On a GPU, these samplers run 23 to 44 times faster than their `jax.random` counterparts on a benchmark, and although approximate, they stay close to the exact distributions: for every parameter value tested, quantiles of their draws differ from those of a reference sampler by at most 1.1% of a standard deviation on average, as measured by the Wasserstein-1 distance (Appendix C).

2 POMP Modeling in Pypomp

A POMP model consists of an unobserved Markov process $\{X(t), t_0 \le t \le t_N\}$ and observations $Y_{1:N}$ made at times $t_{1:N}$, with observed values $y^*_{1:N}$. The initial state is defined at time t_0, $X_n := X(t_n)$, and Y_n is independent of Y_m, X_m, $m \neq n$, conditional on X_n. Plug-and-play methods need only a simulator for the initial state, $X_0 \sim f_{X_0}(x_0;\theta)$, a simulator for the state transition, $X_{n+1} \sim f_{X_{n+1}|X_n}(x_{n+1}|x_n;\theta)$, and the measurement density $f_{Y_n|X_n}(y_n|x_n;\theta)$, where θ denotes the model parameters. In Pypomp, the user passes the data and JAX-compatible programmatic descriptions of these functions to build a model which is a `Pomp` class object. For P&P methods, it suffices to specify an initial value for the latent dynamic system (known as `rinit`), a latent process simulator (`rproc`), a measurement density simulator (`rmeas`) and a measurement density evaluator (`dmeas`). These basic model components are user-specified, and when supplied as arguments in a call to the constructor function, `Pomp`, they become accessible to methods defined for the `Pomp` class. Tutorials at `https://github.com/pypomp/tutorials` provide guidance on building user-specifed models. Demonstration models already included in Pypomp include a model for historical monthly cholera mortality in Dhaka, Bangladesh (`pypomp.models.dacca`), historical weekly measles cases in London, UK (`pypomp.models.measles.pomp`), daily closing price for the

S&P 500 index (`pypomp.models.spx`), and a panel model for historical measles cases in up to 1422 urban and rural districts in England and Wales (`pypomp.models.measles.panel_pomp`).

The following code uses the stochastic volatility example model in Pypomp to demonstrate a search for the maximum likelihood estimate with IF2 (`mif`) followed by refinement using stochastic gradient descent (`train`) and then evaluation of the resulting log-likelihood using a particle filter (`pfilter`). Each method call stores its log-likelihood estimates, parameter traces, execution time, and arguments in the `results_history` attribute of the `Pomp` object, enabling methods to be chained sequentially and providing diagnostic records that can be analyzed by helper methods such as `traces`.

```
import pypomp as pp, jax
spx_obj = pp.models.spx()
spx_obj.mif(J=1000, M=10, key=jax.random.key(1),
    rw_sd = pp.RWSigma(
        sigmas={"mu": 0.02, "kappa": 0.02, "theta": 0.02,
            "xi": 0.02, "rho": 0.02, "V_0": 0.1},
        init_names=["V_0"],
    ).geometric_cooling(a=0.5))
spx_obj.train(J=1000, M=10, optimizer=pp.Adam(),
    eta = pp.LearningRate(
        {name: 0.001 for name in spx_obj.canonical_param_names}
    ).geometric_decay(decay_rate=0.5, M=10))
spx_obj.pfilter(J=1000, reps=10)
```

A collection of POMP models (each referred to as a *unit*) is called a PanelPOMP (Bretó et al., 2020). Typically, some parameters are shared between units and others are unit-specific. Pypomp supports the marginalized panel iterated filtering (MPIF) algorithm (Wheeler et al., 2025) for models of class `PanelPomp`, and DMOP extends directly to PanelPOMP models.

3 Performance

Two performance questions are: (i) What is the benefit of using GPU hardware versus CPU? (ii) How does Pypomp compare in speed to altnernative GPU-based particle filter libraries? To address (i), we compare Pypomp on a GPU with pomp on a CPU because that informs a natural decision for scientists considering Pypomp as a data analysis tool. Model code in pomp is usually written via compiled snippets of C code, and so pomp has comparable speed to other high-performance CPU-based packages for POMP models (FitzJohn et al., 2020). A comparison to address (ii) is given in Appendix[EI: ???? TODO. I SUPPOSE WE WANT THIS IN AN APPENDIX FOR SPACE REASONS.]

Table 1 benchmarks Pypomp 1.0.5--1.0.8 against `pomp` 6.3 and `panelPomp` 1.7.0.0 on four models with different characteristics. The models and additional experimental details are described in Appendix A, with full code at `https://github.com/pypomp/quant`. Relative efficiency varies with the number of particles, with the highest efficiency being obtained when the GPU cores are saturated (Appendix B). The GPU times for the measles models rely on the GPU-adapted samplers in `pypomp.random`: with standard algorithms, such as

Table 1: Wall-clock time in seconds for 36 IF2 searches. Ratio is pomp time divided by Pypomp time, and CPU cores matched is Ratio × 36. Peak cores matched is the highest throughput in terms of CPU cores matched in GPU scaling tests with additional particles (see Appendix B).

Model	`pomp` (36 cores)	Pypomp (1 GPU)	Ratio	CPU cores matched	Peak cores matched
Cholera, Dhaka	1,609	28.3	57×	2,045	3,239
Measles, London	2,923	32.2	91×	3,264	5,583
Measles, 4-town panel	8,262	98.2	84×	3,030	3,947
S&P 500	703	34.8	20×	727	8,659

`jax.random` samplers, IF2 for the London measles model took 46 times longer (Appendix C). Fitting the measles model of He et al. (2010) to all 1,422 towns and cities for 1950–1963 took one GPU 38 hours, including likelihood evaluation, versus an estimated 4.8 months for R on a 36-core node, extrapolating linearly from `panelPomp`'s cost on the 4-town panel (Appendix D).

4 Discussion

The computational cost of PF has hindered likelihood evaluation and statistical inference for POMP models and therefore obstructed scientific investigations. Pypomp's flexibility in model specification, its speed at PF, and the availability of DPF, will allow practitioners to fit new models more quickly, facilitating investigations of model structure. Some models and datasets that were previously considered impractical for likelihood-based inference (Becker and Grenfell, 2017) will become feasible, as demonstrated in Appendix D. Although Pypomp can run on CPU hardware, it gains capability from access to GPUs that are becoming increasingly widely available.

The source code for Pypomp is openly available at `https://github.com/pypomp/pypomp`. The project is run by the Pypomp organization (https://github.com/pypomp) under democratic governance rules that welcomes new developers.

References

Christophe Andrieu, Arnaud Doucet, and Roman Holenstein. Particle Markov chain Monte Carlo methods. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 72(3):269–342, 2010. doi: 10.1111/j.1467-9868.2009.00736.x.

M. Sanjeev Arulampalam, Simon Maskell, Neil Gordon, and Tim Clapp. A tutorial on particle filters for online nonlinear/non-Gaussian Bayesian tracking. *IEEE Transactions on Signal Processing*, 50(2):174–188, 2002. doi: 10.1109/78.978374.

Marie Auger-Méthé, Ken Newman, Diana Cole, Fanny Empacher, Rowenna Gryba, Aaron A. King, Vianey Leos-Barajas, Joanna Mills Flemming, Anders Nielsen, Giovanni

Petris, and Len Thomas. A guide to state-space modeling of ecological time series. *Ecological Monographs*, 91(4):e01470, 2021. doi: 10.1002/ecm.1470.

Alexander D. Becker and Bryan T. Grenfell. tsiR: An R package for time-series Susceptible-Infected-Recovered models of epidemics. *PLOS ONE*, 12(9):e0185528, 2017. doi: 10.1371/journal.pone.0185528.

Julie C. Blackwood, Daniel G. Streicker, Sonia Altizer, and Pejman Rohani. Resolving the roles of immunity, pathogenesis, and immigration for rabies persistence in vampire bats. *Proceedings of the National Academy of Sciences of the United States of America*, 110 (51):20837–20842, 2013. doi: 10.1073/pnas.1308817110.

James Bradbury, Roy Frostig, Peter Hawkins, Matthew James Johnson, Chris Leary, Dougal Maclaurin, George Necula, Adam Paszke, ake VanderPlas, Skye Wanderman-Milne, and Qiao Zhang. JAX: Composable transformations of Python+NumPy programs, 2018.

John-Joseph Brady, Benjamin Cox, Yunpeng Li, and Víctor Elvira. PyDPF: A Python package for differentiable particle filtering, 2026.

Carles Bretó. On idiosyncratic stochasticity of financial leverage effects. *Statistics & Probability Letters*, 91:20–26, 2014. doi: 10.1016/j.spl.2014.04.003.

Carles Bretó, Daihai He, Edward L. Ionides, and Aaron A. King. Time series analysis via mechanistic models. *The Annals of Applied Statistics*, 3(1):319–348, 2009. doi: 10.1214/08-AOAS201.

Carles Bretó, Edward L. Ionides, and Aaron A. King. Panel Data Analysis via Mechanistic Models. *Journal of the American Statistical Association*, 115(531):1178–1188, 2020. doi: 10.1080/01621459.2019.1604367.

Nicolas Chopin and Omiros Papaspiliopoulos. *An Introduction to Sequential Monte Carlo*. Springer, 2020. doi: 10.1007/978-3-030-47845-2.

Sam Duffield. cuthbert: State-space model inference with JAX, 2026. URL `https://pypi.org/project/cuthbert/`.

Richard G. FitzJohn, Edward S. Knock, Lilith K. Whittles, Pablo N. Perez-Guzman, Sangeeta Bhatia, Fernando Guntoro, Oliver J. Watson, Charles Whittaker, Neil M. Ferguson, Anne Cori, Marc Baguelin, and John A. Lees. Reproducible parallel inference and simulation of stochastic state space models using odin, dust, and mcstate. *Wellcome Open Research*, 5, 2020. doi: 10.12688/wellcomeopenres.16466.2.

Spencer J. Fox, Michael Lachmann, Mauricio Tec, Remy Pasco, Spencer Woody, Zhanwei Du, Xutong Wang, Tanyi A. Ingle, E. Javan, Maytal Dahan, Kelly Gaither, Mark E. Escott, Stephen I. Adler, S. Claiborne Johnston, James G. Scott, and Lauren A. Meyers. Real-time pandemic surveillance using hospital admissions and mobility data. *Proceedings of the National Academy of Sciences*, 119(7):e2111870119, 2022. doi: 10.1073/pnas.2111870119.

Michael B. Giles. Algorithm 955: Approximation of the Inverse Poisson Cumulative Distribution Function. *ACM Transactions on Mathematical Software*, 42(1):1–22, 2016. doi: 10.1145/2699466.

Michael B. Giles and Casper Beentjes. Approximation of an inverse of the incomplete Beta Function. In Kevin Buzzard, Alicia Dickenstein, Bettina Eick, Anton Leykin, and Yue Ren, editors, *Mathematical Software – ICMS 2024*, volume 14749, pages 207–214. Springer Nature Switzerland, Cham, 2024. doi: 10.1007/978-3-031-64529-7_22. Series Title: Lecture Notes in Computer Science.

Daihai He, Edward L. Ionides, and Aaron A. King. Plug-and-play inference for disease dynamics: Measles in large and small populations as a case study. *Journal of The Royal Society Interface*, 7(43):271–283, 2010. doi: 10.1098/rsif.2009.0151.

Edward L. Ionides, Dao Nguyen, Yves Atchadé, Stilian Stoev, and Aaron A. King. Inference for dynamic and latent variable models via iterated, perturbed Bayes maps. *Proceedings of the National Academy of Sciences*, 112(3):719–724, 2015. doi: 10.1073/pnas.1410597112.

Edward L. Ionides, Carles Breto, Joonha Park, R. Alexander Smith, and Aaron A. King. Monte Carlo profile confidence intervals for dynamic systems. *Journal of The Royal Society Interface*, 14(132):20170126, 2017. doi: 10.1098/rsif.2017.0126.

Rico Jonschkowski, Divyam Rastogi, and Oliver Brock. Differentiable particle filters: End-to-end learning with algorithmic priors. *Robotics: Science and Systems XIV*, 2018. doi: 10.48550/arXiv.1805.11122.

Santeri Karppinen, Sumeetpal S Singh, and Matti Vihola. Conditional particle filters with bridge backward sampling. *Journal of Computational and Graphical Statistics*, 33(2): 364–378, 2024. doi: 10.1080/10618600.2023.2231514.

Aaron A. King, Edward L. Ionides, Mercedes Pascual, and Menno J. Bouma. Inapparent infections and cholera dynamics. *Nature*, 454(7206):877–880, 2008. doi: 10.1038/nature07084.

Aaron A. King, Dao Nguyen, and Edward L. Ionides. Statistical Inference for Partially Observed Markov Processes via the *R* Package **pomp**. *Journal of Statistical Software*, 69(12), 2016. doi: 10.18637/jss.v069.i12.

Hannah Korevaar, C. Jessica Metcalf, and Bryan T. Grenfell. Structure, space and size: Competing drivers of variation in urban and rural measles transmission. *Journal of The Royal Society Interface*, 17(168):20200010, 2020. doi: 10.1098/rsif.2020.0010.

Scott W. Linderman, Peter Chang, Giles Harper-Donnelly, Aleyna Kara, Xinglong Li, Gerardo Duran-Martin, and Kevin Murphy. Dynamax: A python package for probabilistic state space modeling with JAX. *Journal of Open Source Software*, 10(108):7069, 2025. doi: 10.21105/joss.07069.

Martin Lysy, Pranav Subramani, Jonathan Ramkissoon, Mohan Wu, Michelle Ko, and Kanika Choptra. Pfjax: Particle filtering in JAX, 2022. URL `https://github.com/mlysy/pfjax`.

John A. Jr. Marino, Scott D. Peacor, David B. Bunnell, Henry A. Vanderploeg, Steven A. Pothoven, Ashley K. Elgin, James R. Bence, Jing Jiao, and Edward L. Ionides. Evaluating consumptive and nonconsumptive predator effects on prey density using field time-series data. *Ecology*, 100(3):e02583, 2019. doi: 10.1002/ecy.2583.

Paul Marjoram, John Molitor, Vincent Plagnol, and Simon Tavaré. Markov chain Monte Carlo without likelihoods. *Proceedings of the National Academy of Sciences*, 100(26): 15324–15328, 2003. doi: 10.1073/pnas.0306899100.

Matthew S. Mietchen, Erin Clancey, Corrin McMichael, and Eric T. Lofgren. Estimating SARS-CoV-2 transmission parameters between coinciding outbreaks in a university population and the surrounding community, 2024.

N. Ning, E. L. Ionides, and Y. Ritov. Scalable Monte Carlo inference and rescaled local asymptotic normality. *Bernoulli*, 27:2532–2555, 2021. doi: 10.3150/20-BEJ1321.

G. O. Roberts and O. Stramer. On inference for partially observed nonlinear diffusion models using the Metropolis-Hastings algorithm. *Biometrika*, 88:603–621, 2001. doi: 10.1093/biomet/88.3.603.

Kevin Tan, Aaron J. Abkemeier, Kunyang He, Giles Hooker, and Edward L. Ionides. Automatic differentiation for inference on partially observed Markov process models. *arXiv:2407.03085v2*, 2026. doi: 10.48550/arXiv.2407.03085.

Nico M. Temme. Asymptotic inversion of incomplete gamma functions. *Mathematics of Computation*, 58(198):755–764, 1992. doi: 10.1090/S0025-5718-1992-1122079-8.

Li Wen, Yi Yin, Qiong Li, Zhihang Peng, and Daihai He. Modeling the co-circulation of influenza and COVID-19 in Hong Kong, China. *Advances in Continuous and Discrete Models*, 2024(32(2024)):1–9, 2024. doi: 10.1186/s13662-024-03830-7.

Jesse Wheeler and Edward L Ionides. Revisiting inference for ARMA models: Improved fits and superior confidence intervals. *PloS One*, 20(10):e0333993, 2025. doi: 10.1371/journal.pone.0333993.

Jesse Wheeler, Aaron J. Abkemeier, and Edward L. Ionides. Iterating marginalized Bayes maps for likelihood maximization with application to nonlinear panel models. *arXiv:2511.17438*, 2025. doi: 10.48550/arXiv.2511.17438.

Appendix A. Full benchmark results

The four benchmark models stress different parts of the particle filter, which helps locate performance bottlenecks.

Cholera, Dhaka. The cholera model of King et al. (2008) has 600 monthly observations with 20 Euler steps between each, and each step updates a stochastic differential equation with one normal draw. The process model is cheap, but there are many steps between resamples, so this model isolates the per-step overhead inside the process model.

Measles, London. The measles model of He et al. (2010) has weekly case reports from 1950 to 1963 and takes a daily Euler step, interpolating covariates and drawing binomial, gamma, and Poisson random variables at every step. It has the most expensive process model of the four, so it isolates how the GPU handles heavy computation per particle, and it relies on the approximate samplers in `pypomp.random` (Appendix C).

Measles, 4-town panel. The panel model fits the same measles model to London, Halesworth, Hastings, and Cardiff jointly, with some parameters shared across towns and the rest specific to each town. MPIF (Wheeler et al., 2025) filters the towns sequentially within each iteration, so an iteration costs about four single-town iterations with the same number of particles in flight. This model tests the panel methods used in Appendix D. Its ratio of R [EI: HERE AND ELSEWEHRE: THIS IS CONFUSING, BECAUSE R TIME IS pomp AND GPU TIME IS Pypomp. MAYBE JUST SAY pomp OR DEFINE AND USE R-pomp?] time to GPU time is lower than for the London model because R runs it relatively fast, likely because the random variate generators in R can take a cheaper route for the smaller parameter values of the smaller towns.

S&P 500. The stochastic volatility model of Bretó (2014) has 3,523 daily log returns and takes one step of its process model per observation, with a single normal draw. Each step between resamples is very cheap, so fixed per-observation costs such as resampling can dominate.

R and Pypomp run the same workload in a given benchmark. Each benchmark runs 36 IF2 searches from the same starting points and then runs replicate particle filters at each search's IF2 estimate, with as many particles per filter as per search. The cholera and both measles benchmarks use 5,000 particles, 100 IF2 iterations, and 36 particle filter replicates, and the S&P 500 benchmark uses 1,000 particles, 200 iterations, and 24 replicates.

Table 2 extends Table 1 with particle filter times and Pypomp on the same 36-core CPU node as R, which splits the GPU's advantage into a software part and a hardware part. On the same CPU node, Pypomp is 2.6 to 5.6 times faster than R, and moving Pypomp to the GPU makes it another 1.4 to 33 times faster. These workloads are sized for R and too small to saturate the GPU, so the hardware and total ratios understate what the GPU can do. At its peak throughput in the scaling tests (Appendix B), one GPU runs IF2 90 to 241 times as fast as R on the 36-core node.

Appendix B. GPU scaling tests

The benchmarks in Table 1 and Table 2 use the R configuration, which does not give the GPU enough work to reach its highest throughput for every model. The scaling tests in quant ran IF2 on one RTX PRO 6000 Blackwell Server Edition GPU with more particles in flight, added in two ways: more searches of 1,000 particles each, or more particles for each of a fixed number of searches. Each configuration ran in its own process and ran a short IF2 twice, so that the second run excludes JIT compilation. We project each configuration's time to the benchmark's number of IF2 iterations, including compilation, and divide the benchmark's R time by it to get the number of CPU cores running R that one GPU matches. This assumes that R's time grows linearly with the number of particles and searches. The highest of these values is the Peak column of Table 1. JAX's memory preallocation was turned off, so the peak GPU memory is the memory each configuration used. No scaling

Table 2: Wall-clock time in seconds for IF2 and for particle filter (PF) evaluation, with the workloads described above. Pypomp IF2 and PF times include JIT compilation. *`panelPomp` for the 4-town panel. Software is the `pomp` time divided by the Pypomp time on the same 36 cores, hardware is the Pypomp CPU time divided by the Pypomp GPU time, and total is their product, the `pomp` time divided by the Pypomp GPU time. These workloads do not saturate the GPU, so the hardware and total ratios understate its throughput; peak cores matched in Table 1 gives the GPU's highest measured IF2 throughput. The 4-town panel has no Pypomp CPU run.

		`pomp`*	Pypomp		Ratio		
Model	Task	36 cores	36 cores	1 GPU	Software	Hardware	Total
Cholera, Dhaka	IF2	1,609	390	28.3	4.1×	14×	57×
	PF	380	67.8	8.2	5.6×	8.3×	46×
Measles, London	IF2	2,923	1,071	32.2	2.7×	33×	91×
	PF	895	302	14.5	3.0×	21×	62×
Measles, 4-town panel	IF2	8,262	—	98.2	—	—	84×
	PF	2,252	—	33.7	—	—	67×
S&P 500	IF2	703	275	34.8	2.6×	7.9×	20×
	PF	36.4	7.4	5.2	4.9×	1.4×	7.0×

test used more than 1.8 GB of GPU memory, the peak for the cholera model with 2.4 million particles.

Figure 1 shows the results. For every model, throughput peaks at between 1.0 and 2.0 million particles and falls or plateaus at larger sizes. The peak is 1.6 times the benchmark for the cholera model, 1.7 times for the London measles model, 1.3 times for the 4-town panel, and 11.9 times for the S&P 500 model. The S&P 500 model gains the most because its benchmark, with cheap steps and few particles, leaves the GPU mostly idle. For this model, adding searches beats adding particles per search at a given total number of particles by an average of 6%. Consequently, there is not a large difference in speed between adding searches or particles per search on the GPU in terms of scaling with total particles.

Appendix C. Random variable samplers

`jax.random` draws Poisson, binomial, and gamma random variables with rejection samplers, which loop a random number of times per draw. This suits GPUs poorly, since threads that run together wait for the slowest one. `pypomp.random` instead provides approximate inverse-CDF samplers, which avoid rejection loops. Table 3 compares the two on 1,000,000 draws per call, with parameters spread across the values where each sampler switches algorithms, averaged over 20 calls. The approximation error is small: over a grid of parameter values, the Wasserstein distance between 1,000,000 draws from each `pypomp.random` sampler and from a reference sampler is at most 1.1% of the reference standard deviation. In the London measles benchmark, the `jax.random` samplers made IF2 46 times slower and particle filtering 36 times slower on the GPU.

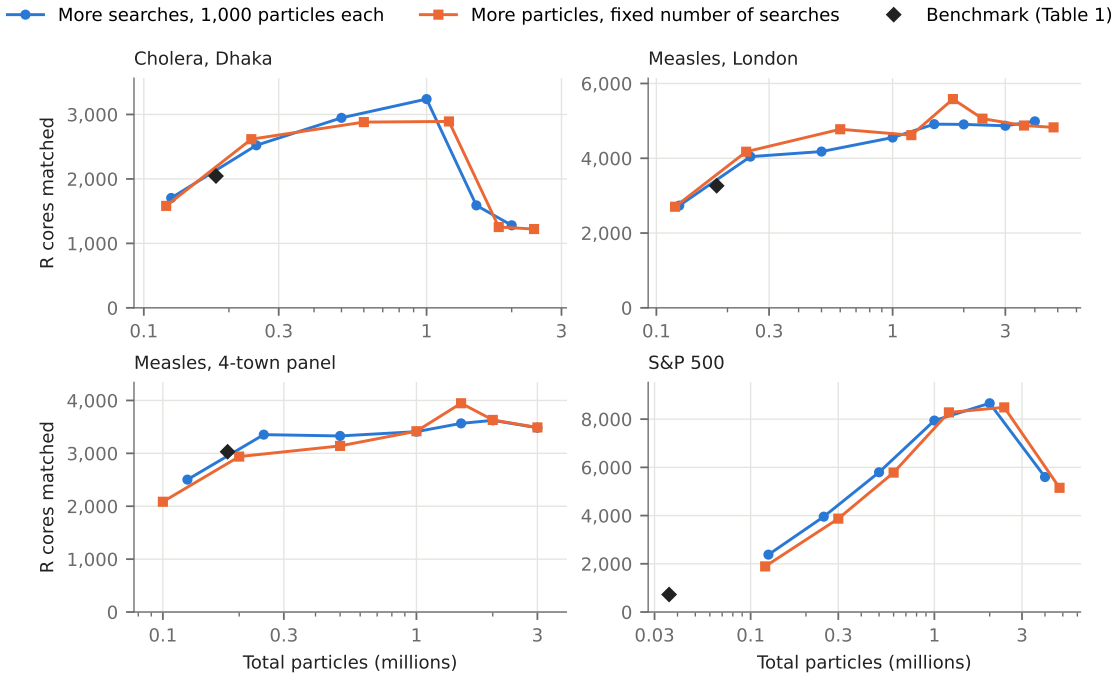


Figure 1: IF2 throughput on one GPU, as the number of CPU cores running R that the GPU matches, against the total number of particles in flight. Blue adds searches of 1,000 particles each; orange adds particles to each of 100 or 120 searches. Each configuration ran 3 or 10 IF2 iterations, projected to the benchmark's number. The x-axes use a log scale.

Table 3: Time in milliseconds for 1,000,000 draws from `pypomp.random` and `jax.random`, on one GPU and on one CPU core, with parameters spread across the values where the samplers switch algorithms. Ratio is the `jax.random` time divided by the `pypomp.random` time. `jax.random` has no negative binomial sampler. Max W_1/σ is the largest Wasserstein distance between `pypomp.random` draws and draws from `jax.random` (or `scipy.stats` for the negative binomial) over a grid of parameter values, relative to the standard deviation.

	GPU			CPU			
Distribution	`pypomp`	`jax`	Ratio	`pypomp`	`jax`	Ratio	Max W_1/σ
Poisson	0.03	1.05	35×	37.17	1,321	36×	0.010
Binomial	0.02	1.07	44×	24.56	1,349	55×	0.010
Gamma	0.03	0.76	23×	118	3,046	26×	0.011
Negative binomial	0.07	—	—	187	—	—	0.009

Table 4: Wall-clock time for fitting the measles model to all towns in England and Wales. Pypomp times were measured on one GPU. R times are estimates for `panelPomp` on a 36-core CPU node, extrapolated from the 4-town panel in Table 2.

Stage	Pypomp (1 GPU)	R `panelPomp`, 36 cores (estimated)
IF2	35.5 hours	136.0 days
PF, all search endpoints	2.5 hours	9.3 days
PF, combined estimate	6 minutes	6.2 hours
Total	38.1 hours	145.5 days

Appendix D. England and Wales measles fit

We ran 36 MPIF searches, each with 10,000 particles for 200 iterations, and evaluated each endpoint with 36 particle filter replicates of 5,000 particles. We then combined the best estimate for each town across searches and evaluated the result with another 36 replicates. Table 4 gives the time for each stage on one GPU with Pypomp 1.0.8.

For comparison, we estimated how long the same computation would take with `panelPomp` on a 36-core CPU node. We measured `panelPomp`'s cost per particle, per IF2 iteration, and per town on the 4-town panel in Table 2, which covers the same period with the same daily time step, and scaled it linearly to the full computation. We assume `panelPomp` fits the towns jointly with MPIF, as Pypomp did, so each search runs on one core and does not parallelize across units – this would be necessary in the general case when using shared parameters. By this estimate, R would need about 4.8 months, 92 times longer than the GPU, and about 14.3 years on a single core, so the GPU delivered the throughput of about 3,300 CPU cores running R. The wall clock time can be reduced further by spreading the replication over multiple GPUs; for example, a run with the same settings but only 12 searches took one GPU 15 hours, so three GPUs running 12 searches each would finish in about 15 hours instead of 38.