

Case Study - Polio in Wisconsin

Aaron A. King Edward L. Ionides Translated in pypomp by Kunyang He

2025-12-24

Table of contents

Introduction	3
Objectives	3
Learning Objectives	3
Covariates	3
Reviewing Covariates in Time Series Analysis	3
What are Covariates?	3
Including Covariates in the General POMP Framework	3
Seasonality in a POMP Model	4
Covariates in pypomp	4
A POMP Model for Polio	4
Polio in Wisconsin	4
Background	4
The Data	5
Loading the Data	5
Visualizing the Data	5
Model Structure	6
Compartment Model	6
Force of Infection	6
Stochastic Force of Infection	7
A pypomp Representation of the POMP Model	7
Building the Model	7
State Variables and Parameters	7
Loading Covariates from R-generated File	8
Fixed Parameters and Starting Values	9
Process Model (rproc)	9
Measurement Model	10
Initial State Distribution	11
Constructing the Pump Object	11
Logistics for the Computations	12
Run Levels	12
Setting Run Levels to Control Computation Time	12

Particle Filtering	12
Likelihood Evaluation	12
Evaluating Likelihood at Starting Parameters	12
Persistence of Polio	13
Simulation Study	13
Investigating Local Persistence	13
Analyzing Persistence	13
Likelihood Maximization	14
Local Search	14
IF2 Random Walk Standard Deviations	14
Running IF2 (Local Search)	14
Evaluating the MIF Results	15
Global Search	15
Setting Up Global Search	15
Running Global Search	15
Analyzing Global Search Results	16
Benchmark Likelihoods	17
Non-Mechanistic Models	17
IID Negative Binomial Benchmark	17
ARMA Benchmark	17
Profile Likelihood	18
Constructing a Profile	18
Profile Over Reporting Rate ρ	18
Summary	18
Key Takeaways	18
What We Learned	18
pypomp API Key Points	18
Exercises	19
Exercise 1: Initial Values	19
Exercise 1: Initial Values	19
Exercise 2: Starting Values	19
Exercise 2: Parameter Estimation Using Randomized Starting Values	19
Acknowledgments and License	19
Acknowledgments	19
References	20

This lesson presents a case study of polio in Wisconsin, demonstrating the use of covariates in **pypomp** and likelihood-based inference for partially observed Markov process models.

Introduction

Objectives

Learning Objectives

1. Demonstrate the use of covariates in **pypomp** to add demographic data (birth rates and total population) and seasonality to an epidemiological model.
2. Show how partially observed Markov process (POMP) models and methods can be used to understand transmission dynamics of polio.
3. Practice maximizing the likelihood for such models. How to set up a global search for a maximum likelihood estimate. How to assess whether a search has been successful.
4. Provide a workflow that can be adapted to related data analysis tasks.

Covariates

Reviewing Covariates in Time Series Analysis

What are Covariates?

- Suppose our time series of primary interest is $y_{1:N}$.
- A *covariate* time series is an additional time series $z_{1:N}$ which is used to help explain $y_{1:N}$.
- When we talk about covariates, it is often implicit that we think of $z_{1:N}$ as a measure of an *external forcing* to the system producing $y_{1:N}$. This means that the process generating the data $z_{1:N}$ affects the process generating $y_{1:N}$, but not vice versa.
- For example, the weather might affect human health, but human health has negligible effect on weather: weather is an external forcing to human health processes.
- When the process leading to $z_{1:N}$ is not external to the system generating it, we must be alert to the possibility of *reverse causation* and *confounding variables*.

Including Covariates in the General POMP Framework

The general POMP modeling framework allows essentially arbitrary modeling of covariates.

Recall that a POMP model is specified by defining, for $n = 1 : N$,

$$\begin{aligned} f_{X_0}(x_0; \theta), \\ f_{X_n|X_{n-1}}(x_n|x_{n-1}; \theta), \\ f_{Y_n|X_n}(y_n|x_n; \theta). \end{aligned}$$

The possibility of a general dependence on n includes the possibility that there is some covariate time series $z_{0:N}$ such that

$$\begin{aligned} f_{X_0}(x_0; \theta) &= f_{X_0}(x_0; \theta, z_0) \\ f_{X_n|X_{n-1}}(x_n|x_{n-1}; \theta) &= f_{X_n|X_{n-1}}(x_n|x_{n-1}; \theta, z_n), \\ f_{Y_n|X_n}(y_n|x_n; \theta) &= f_{Y_n|X_n}(y_n|x_n; \theta, z_n). \end{aligned}$$

Seasonality in a POMP Model

- One specific choice of covariates is to construct $z_{0:N}$ so that it fluctuates periodically, once per year. This allows *seasonality* to enter the POMP model in whatever way is appropriate for the system under investigation.
- All that remains is to hypothesize what is a reasonable way to include covariates for your system, and to fit the resulting model.
- Now we can evaluate and maximize the log-likelihood, we can construct AIC or likelihood ratio tests to see if the covariate helps describe the data.
- This also lets us compare alternative ways the covariates might enter the process model and/or the measurement model.

Covariates in pypomp

- **pypomp** provides facilities for including covariates in a **Pomp** object.
- Covariates entered via the **covars** argument to **Pomp** are automatically interpolated to match observation times and passed to the **rinit**, **rproc**, **dmeas**, and **rmeas** functions.
- We see this in practice in the following epidemiological model, which has population census, birth data and seasonality as covariates.

A POMP Model for Polio

Polio in Wisconsin

Background

- The massive global polio eradication initiative (GPEI) has brought polio from a major global disease to the brink of extinction.
- Finishing this task is proving hard, and improved understanding of polio ecology might assist.
- [Martinez-Bakker et al. \(2015\)](#) investigated this using extensive state level pre-vaccination era data in USA.
- We will follow the approach of [Martinez-Bakker et al. \(2015\)](#) for one state (Wisconsin). In the context of their model, we can quantify seasonality of transmission, the role of the birth rate in explaining the transmission dynamics, and the persistence mechanism of polio.
- [Martinez-Bakker et al. \(2015\)](#) carried out this analysis for all 48 contiguous states and District of Columbia, and their data and code are publicly available.
- The data we study, in **polio_wisconsin.csv**, consist of:
 - **cases**: the monthly reported polio cases
 - **births**: the monthly recorded births
 - **pop**: the annual census
 - **time**: date in years

The Data

Loading the Data

```
import jax
import jax.numpy as jnp
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# For reproducibility
np.random.seed(594709947)

# Load the data
# Note: You can download from https://kingaa.github.io/sbied/polio/polio_wisconsin.csv
data = pd.read_csv("polio_wisconsin.csv", comment='#')
print(data.head())
```

	time	cases	births	pop
0	1931.083333	7	4698	2990000
1	1931.166667	0	4354	2990000
2	1931.250000	7	4836	2990000
3	1931.333333	3	4468	2990000
4	1931.416667	4	4712	2990000

Visualizing the Data


```
fig, axes = plt.subplots(3, 1, figsize=(6, 4), sharex=True)

axes[0].plot(data['time'], data['cases'], 'b-', linewidth=0.5)
axes[0].set_ylabel('Cases')

axes[1].plot(data['time'], data['births'], 'g-', linewidth=0.5)
axes[1].set_ylabel('Births')

axes[2].plot(data['time'], data['pop'], 'r-', linewidth=0.5)
axes[2].set_ylabel('Population')
axes[2].set_xlabel('Year')

plt.tight_layout()
plt.show()
```



Model Structure

Compartment Model

We use the compartment model of [Martinez-Bakker et al. \(2015\)](#):

- Compartments representing susceptible babies in each of six one-month birth cohorts ($S_1^B, \dots, S_6^B$)
- Susceptible older individuals (S^O)
- Infected babies (I^B)
- Infected older individuals (I^O)
- Recovered with lifelong immunity (R)

The state vector of the disease transmission model consists of numbers of individuals in each compartment at each time:

$$X(t) = (S_1^B(t), \dots, S_6^B(t), I^B(t), I^O(t), R(t)).$$

- Babies under six months are modeled as fully protected from symptomatic poliomyelitis.
- Older infections lead to reported cases (usually paralysis) at a rate ρ .

Force of Infection

- Duration of infection is comparable to the one-month reporting aggregation, so a discrete time model may be appropriate.
- [Martinez-Bakker et al. \(2015\)](#) fitted monthly reported cases, May 1932 through January 1953, so we set $t_n = 1932 + (4 + n)/12$ and

$$X_n = X(t_n) = (S_{1,n}^B, \dots, S_{6,n}^B, I_n^B, I_n^O, R_n).$$

- The mean force of infection, in units of yr^{-1} , is modeled as

$$\bar{\lambda}_n = \left(\beta_n \frac{I_n^O + I_n^B}{P_n} + \psi \right)$$

where P_n is census population interpolated to time t_n and seasonality of transmission is modeled as

$$\beta_n = \exp \left\{ \sum_{k=1}^K b_k \xi_k(t_n) \right\},$$

with $\{\xi_k(t), k = 1, \dots, K\}$ a periodic B-spline basis with $K = 6$.

- P_n and $\xi_k(t_n)$ are *covariate time series*.

Stochastic Force of Infection

- The force of infection has a stochastic perturbation,

$$\lambda_n = \bar{\lambda}_n \epsilon_n,$$

where ϵ_n is a Gamma random variable with mean 1 and variance $\sigma_{\text{env}}^2 + \sigma_{\text{dem}}^2 / \bar{\lambda}_n$.

- These two terms capture variation on the environmental and demographic scales, respectively.
- All compartments suffer a mortality rate, set at $\delta = 1/60 \text{ yr}^{-1}$.
- Within each month, all susceptible individuals are modeled as having exposure to constant competing hazards of mortality and polio infection. The chance of remaining in the susceptible population when exposed to these hazards for one month is therefore

$$p_n = \exp \{ -(\delta + \lambda_n)/12 \},$$

with the chance of polio infection being

$$q_n = (1 - p_n) \lambda_n / (\lambda_n + \delta).$$

A pypomp Representation of the POMP Model

Building the Model

State Variables and Parameters

```
# State variable names
statenames = ["SB1", "SB2", "SB3", "SB4", "SB5", "SB6", "IB", "SO", "IO"]

# Initial time
t0 = 1932 + 4/12

# Number of B-spline basis functions for seasonality
K = 6

# We do not explicitly code R, since it is defined implicitly as the
# total population minus the sum of the other compartments.
```

```

# Regular parameters (RPs) - to be estimated
rp_names = ["b1", "b2", "b3", "b4", "b5", "b6",
            "psi", "rho", "tau", "sigma_dem", "sigma_env"]

# Initial value parameters (IVPs) - to be estimated
ivp_names = ["SO_0", "IO_0"]

# Fixed parameters (FPs)
fp_names = ["delta", "SB1_0", "SB2_0", "SB3_0", "SB4_0", "SB5_0", "SB6_0"]

paramnames = rp_names + ivp_names + fp_names

```

Loading Covariates from R-generated File

The covariate file `covar1.csv` is pre-generated from R/pomp to ensure consistency with the original analysis. It contains births (B), population (P), and periodic B-spline basis functions (xi1-xi6).

```

# Load covariates from R-generated file
# This file is generated by pomp to match the original analysis
cov_raw = pd.read_csv("covar1.csv")

# The file has variable names in first column, times as column headers
# We need to transpose it
first_col = cov_raw.columns[0]
covars = (cov_raw.rename(columns={first_col: "var"})
          .set_index("var")
          .T)

# Convert index to float (time values)
covars.index = covars.index.astype(float)
covars.index.name = "time"
covars = covars.sort_index()

print(f"Covariate columns: {list(covars.columns)}")
print(f"Time range: {covars.index[0]:.4f} to {covars.index[-1]:.4f}")
print(covars.head())

```

Covariate columns: ['B', 'xi1', 'xi2', 'xi3', 'xi4', 'xi5', 'xi6', 'P']

Time range: 1931.0833 to 1955.0000

var	B	xi1	xi2	xi3	xi4	xi5	\
time							
1931.083333	4698.0	4.791667e-01	0.479167	0.020833	0.000000e+00	0.000000	
1931.166667	4354.0	1.666667e-01	0.666667	0.166667	1.246144e-33	0.000000	
1931.250000	4836.0	2.083333e-02	0.479167	0.479167	2.083333e-02	0.000000	
1931.333333	4468.0	1.246112e-33	0.166667	0.666667	1.666667e-01	0.000000	
1931.416667	4712.0	0.000000e+00	0.020833	0.479167	4.791667e-01	0.020833	

var	xi6	P
time		
1931.083333	0.020833	2.993754e+06
1931.166667	0.000000	2.996197e+06
1931.250000	0.000000	2.998629e+06
1931.333333	0.000000	3.001045e+06
1931.416667	0.000000	3.003439e+06

Fixed Parameters and Starting Values

```
# Find initial births for SB1_0 through SB6_0 from covariate table
def get_initial_births(covars_df, t0_val):
    """Extract initial birth values from covariate table."""
    idx0 = np.argmin(np.abs(covars_df.index.values - t0_val))
    B_series = covars_df["B"].values
    out = []
    for k_back in range(0, 6):
        j = max(0, idx0 - k_back)
        out.append(float(B_series[j]))
    return out

SB0_list = get_initial_births(covars, t0)

fixed_params = {
    'delta': 1/60, # Mortality rate (per year)
    'SB1_0': SB0_list[0],
    'SB2_0': SB0_list[1],
    'SB3_0': SB0_list[2],
    'SB4_0': SB0_list[3],
    'SB5_0': SB0_list[4],
    'SB6_0': SB0_list[5],
}

print("Fixed parameters:")
for k, v in fixed_params.items():
    print(f" {k}: {v}")
```

Fixed parameters:

```
delta: 0.016666666666666666
SB1_0: 4069.0
SB2_0: 4565.0
SB3_0: 4410.0
SB4_0: 4616.0
SB5_0: 4305.0
SB6_0: 4032.0
```

```
# Starting values for parameters (based on informal model exploration)
params_guess = {
    'b1': 3.0, 'b2': 0.0, 'b3': 1.5, 'b4': 6.0, 'b5': 5.0, 'b6': 3.0,
    'psi': 0.002, 'rho': 0.01, 'tau': 0.001,
    'sigma_dem': 0.04, 'sigma_env': 0.5,
    'S0_0': 0.12, 'I0_0': 0.001,
    **fixed_params
}
```

Process Model (rproc)

```
def rproc(X_, theta_, key, covars, t, dt):
    """
    Process model for polio transmission.

    This follows the Martinez-Bakker et al. (2015) model structure.
    """
    # Extract parameters
    b1, b2, b3, b4, b5, b6 = theta_['b1'], theta_['b2'], theta_['b3'], \
        theta_['b4'], theta_['b5'], theta_['b6']
    psi = theta_['psi']
    sigma_dem = theta_['sigma_dem']
    sigma_env = theta_['sigma_env']
    delta = theta_['delta']
```

```

# Extract states
SB1, SB2, SB3, SB4, SB5, SB6 = X_['SB1'], X_['SB2'], X_['SB3'], \
                                X_['SB4'], X_['SB5'], X_['SB6']
IB, S0, IO = X_['IB'], X_['S0'], X_['IO']

# Extract covariates (from R-generated covar1.csv)
B = covars['B']
P = covars['P']
xi1, xi2, xi3 = covars['xi1'], covars['xi2'], covars['xi3']
xi4, xi5, xi6 = covars['xi4'], covars['xi5'], covars['xi6']

# Compute beta from B-spline basis
beta = jnp.exp(b1*xi1 + b2*xi2 + b3*xi3 + b4*xi4 + b5*xi5 + b6*xi6)

# Mean force of infection
lambda_bar = beta * (IO + IB) / P + psi

# Stochastic perturbation (Gamma noise)
var_epsilon = sigma_dem**2 / jnp.maximum(lambda_bar, 1e-12) + sigma_env**2

# Sample from Gamma distribution with mean 1 and variance var_epsilon
key, subkey = jax.random.split(key)
epsilon = jax.lax.cond(
    var_epsilon < 1e-6,
    lambda _: 1.0,
    lambda _: pp.random.fast_gamma(subkey, 1.0/var_epsilon) * var_epsilon,
    operand=None
)

lambda_n = lambda_bar * epsilon

# Transition probabilities
p = jnp.exp(-(delta + lambda_n) / 12.0)
q = (1.0 - p) * lambda_n / (delta + lambda_n)

# Update states following Martinez-Bakker model
new_SB1 = B
new_SB2 = new_SB1 * p
new_SB3 = new_SB2 * p
new_SB4 = new_SB3 * p
new_SB5 = new_SB4 * p
new_SB6 = new_SB5 * p
new_S0 = (new_SB6 + S0) * p
new_IB = (new_SB1 + new_SB2 + new_SB3 + new_SB4 + new_SB5 + new_SB6) * q
new_IO = new_S0 * q

return {
    'SB1': new_SB1, 'SB2': new_SB2, 'SB3': new_SB3,
    'SB4': new_SB4, 'SB5': new_SB5, 'SB6': new_SB6,
    'IB': new_IB, 'S0': new_S0, 'IO': new_IO
}

```

Measurement Model

```

def dmeas(Y_, X_, theta_, covars, t):
    """
    Measurement density: discretized truncated normal with numerical stability.
    """
    rho = theta_['rho']
    tau = theta_['tau']
    IO = X_['IO']
    cases = Y_['cases']

    mu = rho * IO
    # Add base variance (1.0) to prevent distribution from being too narrow
    sd = jnp.sqrt((tau * IO)**2 + jnp.maximum(mu, 0.0) + 1.0)

```

```

def _cdf(z):
    return 0.5 * (1.0 + jax.scipy.special.erf((z - mu) / (jnp.sqrt(2.0) * sd)))

tol = 1e-18
prob = jax.lax.cond(
    cases > 0.0,
    lambda _: _cdf(cases + 0.5) - _cdf(cases - 0.5) + tol,
    lambda _: _cdf(0.5) + tol,
    operand=None
)

# Ensure probability has a lower bound to prevent log(0)
prob = jnp.maximum(prob, 1e-300)

return jnp.log(prob)

```

```

def rmeas(X_, theta_, key, covars, t):
    """
    Measurement simulator: sample observations given state.

    IMPORTANT: rmeas must return a JAX array with shape (ydim,).
    """
    rho = theta_['rho']
    tau = theta_['tau']
    IO = X_['IO']

    mu = rho * IO
    sd = jnp.sqrt((tau * IO)**2 + jnp.maximum(mu, 0.0))

    z = jax.random.normal(key)
    cases_cont = mu + sd * z
    cases = jnp.where(cases_cont > 0.0, jnp.round(cases_cont), 0.0)

    return jnp.array([cases])

```

Initial State Distribution

```

def rinit(theta_, key, covars, t0):
    """
    Initial state distribution.
    """
    P = covars['P']

    return {
        'SB1': theta_['SB1_0'],
        'SB2': theta_['SB2_0'],
        'SB3': theta_['SB3_0'],
        'SB4': theta_['SB4_0'],
        'SB5': theta_['SB5_0'],
        'SB6': theta_['SB6_0'],
        'IB': 0.0,
        'IO': theta_['IO_0'] * P,
        'SO': theta_['SO_0'] * P
    }

```

Constructing the Pomp Object

```

import pypomp as pp

# Filter data to the analysis period
data_filtered = data[(data['time'] > t0 + 0.01) & (data['time'] < 1953 + 1/12 + 0.01)].copy()

# Create observations DataFrame
ys = data_filtered[['cases']].copy()
ys.index = data_filtered['time']

```

```
# Filter covariates to match observation times
ys = ys.loc[ys.index.intersection(covars.index)]

# Create pomp object
polio = pp.Pomp(
    ys=ys,
    theta=params_guess,
    rinit=rinit,
    rproc=rproc,
    dmeas=dmeas,
    rmeas=rmeas,
    covars=covars,
    statenames=statenames,
    t0=t0,
    nstep=1
)

print(f"Pomp object created")
print(f"Number of observations: {len(ys)}")
print(f"Time range: {ys.index[0]:.3f} to {ys.index[-1]:.3f}")
```

Pomp object created
 Number of observations: 249
 Time range: 1932.417 to 1953.083

Logistics for the Computations

Run Levels

Setting Run Levels to Control Computation Time

```
run_level = 2

Np = [100, 1000, 5000][run_level - 1]
Nmif = [10, 100, 200][run_level - 1]
Nreps_eval = [2, 10, 20][run_level - 1]
Nreps_local = [10, 20, 40][run_level - 1]
Nreps_global = [10, 20, 100][run_level - 1]
Nsim = [50, 100, 500][run_level - 1]

print(f"Run level: {run_level}")
print(f"Np={Np}, Nmif={Nmif}, Nreps_eval={Nreps_eval}")
```

Run level: 2
 Np=1000, Nmif=100, Nreps_eval=10

Particle Filtering

Likelihood Evaluation

Evaluating Likelihood at Starting Parameters

```
params_rpomp = {
    'b1': 3.0, 'b2': 0.0, 'b3': 1.5, 'b4': 6.0, 'b5': 5.0, 'b6': 3.0,
    'psi': 0.002, 'rho': 0.01, 'tau': 0.001,
    'sigma_dem': 0.04, 'sigma_env': 0.5,
    'SQ_0': 0.12, 'IQ_0': 0.001,
    'delta': 1/60, # 0.01666667
    'SB1_0': 4069.0, 'SB2_0': 4565.0, 'SB3_0': 4410.0,
    'SB4_0': 4616.0, 'SB5_0': 4305.0, 'SB6_0': 4032.0
```

```

}

polio.theta = params_rpomp

key = jax.random.key(3899882)
polio.pfilter(J=Np, key=key, reps=20, thresh=0)

pf_result = polio.results_history[-1]
logliks = pf_result.logLiks.values.flatten()

L1 = pp.maths.logmeanexp(logliks)
L1_se = pp.maths.logmeanexp_se(logliks)

print(f"pypomp Log-likelihood: {L1:.2f} (SE: {L1_se:.2f})")

```

pypomp Log-likelihood: -820.77 (SE: 0.72)

Persistence of Polio

Simulation Study

Investigating Local Persistence

```

key = jax.random.key(1643079359)
X_sims, Y_sims = polio.simulate(key=key, nsim=Nsim)

print(f"Simulated {Nsim} trajectories")

```

Simulated 100 trajectories

Analyzing Persistence


```

fig, ax = plt.subplots(figsize=(5, 3))
ax.plot(ys.index, ys['cases'], 'k-', linewidth=1, label='Data', alpha=0.8)

for i in range(min(5, Nsim)):
    sim_data = Y_sims[Y_sims['sim'] == i]
    ax.plot(sim_data['time'], sim_data['obs_0'], alpha=0.3, linewidth=0.5)

ax.set_xlabel('Year')
ax.set_ylabel('Cases')
ax.set_title('Simulations from Fitted Model')
ax.legend()
plt.tight_layout()
plt.show()

```



Likelihood Maximization

Local Search

IF2 Random Walk Standard Deviations

```
rw_sd = pp.RWSigma(
    sigmas={
        'b1': 0.02, 'b2': 0.02, 'b3': 0.02,
        'b4': 0.02, 'b5': 0.02, 'b6': 0.02,
        'psi': 0.02, 'rho': 0.02, 'tau': 0.02,
        'sigma_dem': 0.02, 'sigma_env': 0.02,
        'I0_0': 0.2, 'S0_0': 0.2,
        'delta': 0.0,
        'SB1_0': 0.0, 'SB2_0': 0.0, 'SB3_0': 0.0,
        'SB4_0': 0.0, 'SB5_0': 0.0, 'SB6_0': 0.0
    },
    init_names=['I0_0', 'S0_0']
)
```

Running IF2 (Local Search)

```
key = jax.random.key(942098028)
np.random.seed(12345)
n_starts = min(5, Nreps_local)
theta_list = [params_guess.copy() for _ in range(n_starts)]

polio.mif(J=Np, M=Nmif, key=key, rw_sd=rw_sd, a=0.5, theta=theta_list)

mif_result = polio.results_history[-1]
print(f"MIF completed")

# MIF result stores final parameters, not logLiks directly
# We need to run pfilter to evaluate the likelihood
print(f"Number of parameter sets after MIF: {len(polio.theta)}")
```

MIF completed

Number of parameter sets after MIF: 5

Evaluating the MIF Results

```
# Run pfilter to evaluate log-likelihoods at the MIF-optimized parameters
polio.pfilter(J=Np, key=jax.random.key(12345), reps=Nreps_eval, thresh=0)

pf_result = polio.results_history[-1]
logliks = pf_result.logLiks.values

# Get mean loglik for each parameter set
mean_logliks = logliks.mean(axis=1)
best_idx = np.argmax(mean_logliks)

ll_mif = pp.maths.logmeanexp(logliks[best_idx, :])
ll_mif_se = pp.maths.logmeanexp_se(logliks[best_idx, :])

print(f"Log-likelihood after local search: {ll_mif:.2f} (SE: {ll_mif_se:.2f})")
print(f"Best parameter set index: {best_idx}")
```

Log-likelihood after local search: -1069.38 (SE: 6.82)

Best parameter set index: 3

Global Search

Setting Up Global Search

```
box = {
    'b1': (-2, 8), 'b2': (-2, 8), 'b3': (-2, 8),
    'b4': (-2, 8), 'b5': (-2, 8), 'b6': (-2, 8),
    'psi': (0, 0.1), 'rho': (0, 0.1), 'tau': (0, 0.1),
    'sigma_dem': (0, 0.5), 'sigma_env': (0, 1),
    'SO_0': (0, 1), 'IO_0': (0, 0.01)
}
```

Running Global Search

```
def random_params_from_box(box, n, fixed_params):
    params_list = []
    for _ in range(n):
        params = {}
        for name, (lo, hi) in box.items():
            params[name] = np.random.uniform(lo, hi)
        params.update(fixed_params)
        params_list.append(params)
    return params_list

np.random.seed(833102018)
starts = random_params_from_box(box, min(5, Nreps_global), fixed_params)
print(f"Generated {len(starts)} random starting points")
```

Generated 5 random starting points

```
key = jax.random.key(71449038)
polio.theta = starts
polio.mif(J=Np, M=Nmif, key=key, rw_sd=rw_sd, a=0.5)

# Evaluate likelihoods at the optimized parameters
polio.pfilter(J=Np, key=jax.random.key(99999), reps=Nreps_eval, thresh=0)

pf_result = polio.results_history[-1]
logliks = pf_result.logLiks.values

results = []
for i, theta_i in enumerate(polio.theta):
```

```

    ll = logliks[i, :].mean()
    results.append({'logLik': ll, **theta_i})
    print(f"  Search {i+1}: logLik = {ll:.2f}")

results_df = pd.DataFrame(results)
print(f"\nBest log-likelihood: {results_df['logLik'].max():.2f}")

```

```

Search 1: logLik = -1230.87
Search 2: logLik = -1817.36
Search 3: logLik = -1017.09
Search 4: logLik = -1327.21
Search 5: logLik = -1080.13

```

Best log-likelihood: -1017.09

Analyzing Global Search Results

```

fig, axes = plt.subplots(2, 2, figsize=(5, 4))

good_results = results_df[results_df['logLik'] > results_df['logLik'].max() - 20]

axes[0,0].scatter(good_results['rho'], good_results['logLik'], s=10)
axes[0,0].set_xlabel('rho')
axes[0,0].set_ylabel('logLik')


axes[0,1].scatter(good_results['psi'], good_results['logLik'], s=10)
axes[0,1].set_xlabel('psi')
axes[0,1].set_ylabel('logLik')

axes[1,0].scatter(good_results['sigma_dem'], good_results['logLik'], s=10)
axes[1,0].set_xlabel('sigma_dem')
axes[1,0].set_ylabel('logLik')

axes[1,1].scatter(good_results['sigma_env'], good_results['logLik'], s=10)
axes[1,1].set_xlabel('sigma_env')
axes[1,1].set_ylabel('logLik')

plt.tight_layout()
plt.show()

```

Benchmark Likelihoods

Non-Mechanistic Models

IID Negative Binomial Benchmark

```
from scipy.optimize import minimize
from scipy.stats import nbinom

def nb_negloglik(params, cases):
    size = np.exp(params[0])
    prob = 1 / (1 + np.exp(-params[1]))
    return -np.sum(nbinom.logpmf(cases.astype(int), size, prob))

cases = ys['cases'].values
result = minimize(nb_negloglik, [0, -5], args=(cases,))
nb_loglik = -result.fun

print(f"Negative binomial (IID) log-likelihood: {nb_loglik:.1f}")
```

Negative binomial (IID) log-likelihood: -1036.2

ARMA Benchmark

```
from statsmodels.tsa.arima.model import ARIMA

log_y = np.log(cases + 1)
arma_model = ARIMA(log_y, order=(2, 0, 2), seasonal_order=(1, 0, 1, 12))
arma_fit = arma_model.fit()

arma_loglik = arma_fit.llf - np.sum(log_y)

print(f"ARMA(2,0,2)x(1,0,1)_12 log-likelihood: {arma_loglik:.1f}")
```

ARMA(2,0,2)x(1,0,1)_12 log-likelihood: -824.3

Profile Likelihood

Constructing a Profile

Profile Over Reporting Rate ρ

```
profile_pts = [3, 5, 30][run_level - 1]
rho_values = np.linspace(0.01, 0.025, profile_pts)
print(f"Profiling over {profile_pts} values of rho")

rw_sd_profile = pp.RWSigma(
    sigmas={
        'b1': 0.02, 'b2': 0.02, 'b3': 0.02,
        'b4': 0.02, 'b5': 0.02, 'b6': 0.02,
        'psi': 0.02, 'rho': 0.0, # Fixed!
        'tau': 0.02, 'sigma_dem': 0.02, 'sigma_env': 0.02,
        'IO_0': 0.2, 'SO_0': 0.2,
        'delta': 0.0,
        'SB1_0': 0.0, 'SB2_0': 0.0, 'SB3_0': 0.0,
        'SB4_0': 0.0, 'SB5_0': 0.0, 'SB6_0': 0.0
    },
    init_names=['IO_0', 'SO_0']
)

for rho_val in rho_values[:3]:
    print(f" Would profile at rho = {rho_val:.4f}")
```

```
Profiling over 5 values of rho
Would profile at rho = 0.0100
Would profile at rho = 0.0138
Would profile at rho = 0.0175
```

Summary

Key Takeaways

What We Learned

1. **Covariates** can be included naturally in POMP models to incorporate external forcing like birth rates, population, and seasonality.
2. **Seasonal transmission** is modeled using periodic B-spline basis functions, allowing flexible seasonality patterns.
3. **Global search** from diverse starting points is essential to have confidence in MLEs for complex models.
4. **Profile likelihood** provides rigorous uncertainty quantification for individual parameters.
5. **Benchmark comparisons** (IID, ARMA) help interpret whether mechanistic models are performing adequately.

pypomp API Key Points

1. RWSigma requires **ALL** parameters in sigmas dict - use 0.0 for fixed params

2. `rmeas` must return a JAX array with shape `(ydim,)`, e.g., `jnp.array([value])`
3. Use `nstep=1` for discrete-time models
4. Multiple starting points: pass `theta=[list_of_dicts]` to `mif()`
5. After `mif()`, run `pfilter()` to evaluate log-likelihoods
6. Access `pfilter` results via `pf_result.logLiks.values`

Exercises

Exercise 1: Initial Values

Exercise 1: Initial Values

When carrying out parameter estimation for dynamic systems, we need to specify beginning values for both the dynamic system (in the state space) and the parameters (in the parameter space).

Tasks:

1. Discuss issues in specifying and inferring initial conditions, with particular reference to this polio example.
2. Suggest a possible improvement in the treatment of initial conditions here, code it up and make some preliminary assessment of its effectiveness.
3. How will you decide if it is a substantial improvement?

Exercise 2: Starting Values

Exercise 2: Parameter Estimation Using Randomized Starting Values

Think about possible improvements on the assignment of randomized starting values for the parameter estimation searches.

Tasks:

1. Propose and try out a modification of the procedure.
2. Does it make a difference?

Acknowledgments and License

Acknowledgments

- This lesson is prepared for the Simulation-based Inference for Epidemiological Dynamics module at the Summer Institute in Statistics and Modeling in Infectious Diseases (SISMID).
- The materials build on previous versions of this course and related courses.
- Licensed under the Creative Commons Attribution-NonCommercial license (CC BY-NC 4.0). Please share and remix non-commercially, mentioning its origin.

References

Martinez-Bakker, M., King, A. A., and Rohani, P. (2015). Unraveling the transmission ecology of polio. *PLoS Biol*, 13(6):e1002172.